ERA predictor

January 16, 2019

1 Year to year ERA prediction

This notebook analyzes which statistics best predict a pitchers Earned Run Average, (ERA) from one year to the next

```
In [1]: import pandas as pd
        import numpy as np
        import matplotlib.pyplot as plt
        from sklearn.linear_model import LinearRegression
        from sklearn.model_selection import train_test_split
        from sklearn.metrics import mean_squared_error
        from sklearn.preprocessing import StandardScaler
        from sklearn.tree import DecisionTreeRegressor
        %matplotlib inline
```

We first import select pitching data for qualified starting pitchers from www.fangraphs.com in the date range 2002-2017, for an explanation of each statistic see https://www.fangraphs.com/library/pitching/ .

```
In [2]: stats=pd.read_csv('FanGraphs Leaderboard.csv')
```

```
In [3]: stats.head()
```

```
Out[3]:    Season            Name     Team  Age    K/9  BB/9  K/BB   H/9  HR/9  WHIP  \
        0    2015    Zack Greinke  Dodgers   31   8.08  1.62  5.00  5.98  0.57  0.84
        1    2015    Jake Arrieta     Cubs   29   9.28  1.89  4.92  5.90  0.39  0.86
        2    2014 Clayton Kershaw  Dodgers   26  10.85  1.41  7.71  6.31  0.41  0.86
        3    2013 Clayton Kershaw  Dodgers   25   8.85  1.98  4.46  6.25  0.42  0.92
        4    2005   Roger Clemens   Astros   42   7.88  2.64  2.98  6.43  0.47  1.01

             ...     F-Strike%       K%     BB%    Soft%     Med%    Hard%   ERA  xFIP  \
        0    ...        64.1 %   23.7 %   4.7 %   21.7 %   51.5 %   26.8 %  1.66  3.22
        1    ...        60.2 %   27.1 %   5.5 %   22.8 %   55.2 %   22.1 %  1.77  2.61
        2    ...        68.8 %   31.9 %   4.1 %   24.5 %   51.2 %   24.3 %  1.77  2.08
        3    ...        65.1 %   25.6 %   5.7 %   14.4 %   56.7 %   28.9 %  1.83  2.88
        4    ...        60.4 %   22.1 %   7.4 %   16.3 %   60.9 %   22.8 %  1.87  3.31
```

```
  SIERA  playerid
0  3.27      1943
1  2.75      4153
2  2.09      2036
3  2.99      2036
4  3.47       815

[5 rows x 35 columns]
```

One sees that there is a column for each statistic and a row for each pitcher for each season

```
In [4]: stats.columns

Out[4]: Index(['Season', 'Name', 'Team', 'Age', 'K/9', 'BB/9', 'K/BB', 'H/9', 'HR/9',
               'WHIP', 'BABIP', 'GB/FB', 'LD%', 'GB%', 'FB%', 'IFFB%', 'HR/FB',
               'O-Swing%', 'Z-Swing%', 'Swing%', 'O-Contact%', 'Z-Contact%',
               'Contact%', 'Zone%', 'SwStr%', 'F-Strike%', 'K%', 'BB%', 'Soft%',
               'Med%', 'Hard%', 'ERA', 'xFIP', 'SIERA', 'playerid'],
              dtype='object')
```

The included stats are listed above, I have only included rate stats because counting statistics will skew predicitions in favor of pitchers who have pitched more innings in a given season.

Below I reformat the data values in each column so that they are all floats and are suitable for analysis

```
In [41]: for col in stats.columns:
             if type(stats[col][1])==str:
                 stats[col]=stats[col].str.replace('%','',regex=False)


In [42]: numbers=['Age', 'K/9', 'BB/9', 'K/BB', 'H/9', 'HR/9',
                 'WHIP', 'BABIP', 'GB/FB', 'LD%', 'GB%', 'FB%', 'IFFB%', 'HR/FB',
                 'O-Swing%', 'Z-Swing%', 'Swing%', 'O-Contact%', 'Z-Contact%',
                 'Contact%', 'Zone%', 'SwStr%', 'F-Strike%', 'K%', 'BB%', 'Soft%',
                 'Med%', 'Hard%', 'ERA', 'playerid','xFIP','SIERA']
         for col in numbers:
             stats[col]=stats[col].astype(float)
```

In the next three cells I create a new dataframe from the original that includes a column for the given pitchers ERA the following season, and rename the columns accordingly

```
In [5]: predseasons=['stats'+str(i)+str(i+1) for i in range(2002,2017)]
        i=2002
        for season in predseasons:

            globals()[season] = pd.merge(stats[stats['Season']==i],stats[stats['Season']==(i+1)]
            i+=1


In [6]: statsrel=pd.concat([globals()[season] for season in predseasons])
```

```
In [7]: statsrel=statsrel.rename(columns={'ERA_x':'ERA_current_year','ERA_y':'ERA_next_year'})
```

```
In [8]: statsrel.head()
```

```
Out[8]:    Season             Name       Team  Age    K/9  BB/9  K/BB   H/9  HR/9  \
        0    2002  Pedro Martinez    Red Sox   30  10.79  1.81  5.98  6.50  0.59
        1    2002      Derek Lowe    Red Sox   29   5.20  1.97  2.65  6.80  0.49
        2    2002     Greg Maddux     Braves   36   5.33  2.03  2.62  8.76  0.63
        3    2002      Barry Zito  Athletics   24   7.14  3.06  2.33  7.14  0.94
        4    2002   Bartolo Colon      - - -   29   5.75  2.70  2.13  8.45  0.77

           WHIP      ...            K%     BB%    Soft%     Med%    Hard% ERA_current_year   \
        0  0.92      ...        30.4 %   5.1 %   13.8 %   67.6 %   18.7 %             2.26
        1  0.97      ...        14.9 %   5.6 %   13.3 %   71.2 %   15.4 %             2.58
        2  1.20      ...        14.4 %   5.5 %   15.6 %   66.8 %   17.6 %             2.62
        3  1.13      ...        19.4 %   8.3 %   18.7 %   61.4 %   19.9 %             2.75
        4  1.24      ...        15.4 %   7.3 %   17.4 %   61.1 %   21.5 %             2.93

           xFIP SIERA playerid ERA_next_year
        0  2.61  2.43      200          2.22
        1  3.42  3.18      199          4.47
        2  3.61  3.75      104          3.96
        3  4.28  4.13      944          3.30
        4  4.06  4.26      375          3.87

        [5 rows x 36 columns]
```

"features" is a list of statistics that we will use to try to predict a pitchers ERA.

"featureswithestimators" also includes the ERA estimators xFIP (eXpected Fielding Independent Pitching) and SIERA (Skill Interactive ERA), these stats use strikeout, walk and fly ball rates to say what a pitchers ERA 'should' be by removing factors such as team defense that are out of the pitchers control. For more information on these stats see https://www.fangraphs.com/library/pitching/.

```
In [99]: features=['Age', 'K/9', 'BB/9', 'K/BB', 'H/9', 'HR/9',
                 'WHIP', 'BABIP', 'GB/FB', 'LD%', 'GB%', 'FB%', 'IFFB%', 'HR/FB',
                 'O-Swing%', 'Z-Swing%', 'Swing%', 'O-Contact%', 'Z-Contact%',
                 'Contact%', 'Zone%', 'SwStr%', 'F-Strike%', 'K%', 'BB%', 'Soft%',
                 'Med%', 'Hard%', 'ERA_current_year']
```

```
In [100]: featureswithestimators=['Age', 'K/9', 'BB/9', 'K/BB', 'H/9', 'HR/9',
                  'WHIP', 'BABIP', 'GB/FB', 'LD%', 'GB%', 'FB%', 'IFFB%', 'HR/FB',
                  'O-Swing%', 'Z-Swing%', 'Swing%', 'O-Contact%', 'Z-Contact%',
                  'Contact%', 'Zone%', 'SwStr%', 'F-Strike%', 'K%', 'BB%', 'Soft%',
                  'Med%', 'Hard%', 'ERA_current_year','xFIP','SIERA']
```

First lets explore which stats are most correlated with a pitcher's ERA during the same year

```
In [102]: corrdict1={col:np.abs(statsrel[col].corr(statsrel['ERA_current_year'])) for col in fea
```

```
Corr=pd.DataFrame.from_dict(corrdict1, orient='index')
Corr.sort_values(by=0,ascending=False)
```

Out[102]:

	0
ERA_current_year	1.000000
WHIP	0.815086
H/9	0.759634
SIERA	0.638081
xFIP	0.634663
HR/9	0.590251
K%	0.534908
K/BB	0.489126
K/9	0.452299
SwStr%	0.433424
BABIP	0.432780
HR/FB	0.423012
Contact%	0.411598
O-Swing%	0.353715
Z-Contact%	0.338963
Soft%	0.304414
BB/9	0.299132
F-Strike%	0.248818
Swing%	0.233694
Hard%	0.230031
BB%	0.218658
LD%	0.144044
O-Contact%	0.125753
Zone%	0.114846
GB%	0.107033
GB/FB	0.090136
IFFB%	0.080803
Z-Swing%	0.071524
Age	0.064609
FB%	0.059323
Med%	0.039181

We see that WHIP and hits per nine innings have strong correlations to ERA, placing just above xFIP and SIERA as the most strongly correlated variables. Now lets take a look at how these stats are correlated to a pitcher's ERA the following season.

In [103]:
```
corrdict2={col:np.abs(statsrel[col].corr(statsrel['ERA_next_year'])) for col in featur
```

```
Corr=pd.DataFrame.from_dict(corrdict2, orient='index')
Corr.sort_values(by=0,ascending=False)
```

Out[103]:

	0
SIERA	0.453758
xFIP	0.445884
K%	0.424969
K/9	0.406644
ERA_current_year	0.339578
K/BB	0.335974
SwStr%	0.319018
WHIP	0.316859
Contact%	0.309912
H/9	0.302348
Z-Contact%	0.290058
HR/9	0.278060
O-Swing%	0.251271
Soft%	0.170777
F-Strike%	0.160567
HR/FB	0.156215
Swing%	0.145038
Age	0.117504
Med%	0.111272
BB/9	0.106611
Zone%	0.103316
BB%	0.079882
GB/FB	0.056518
GB%	0.054717
O-Contact%	0.054508
Z-Swing%	0.049796
FB%	0.042389
LD%	0.038722
IFFB%	0.016310
Hard%	0.015875
BABIP	0.014046

Here we see a very different picture, SIERA and xFIP have moved to the front of the pack, justifying their use as truer measures of a pitchers talent than ERA alone. We see that WHIP, and H/9 are now only the 8th and 10th most correlated stats respectively.

Besides xFIP and SIERA, it seems that the best raw stats to use to predict a pitchers ERA are the ones involving strikeouts, (K%, K/9,K/BB, SwStr%). Its no surprise then that strikeouts are a major component of both the SIERA and xFIP estimators.

Also of note is that a pitchers walk rate, (BB%) has a very low correlation to his ERA the next year. It seems that walks do not hurt a pitcher very much as long as he maintains high strikeout totals.

Finally, we try to use a combination of these features in a linear regressor to predict a pitchers ERA the following year.

First we build our design matrix and scale our features accordingly.

```
In [77]: X = StandardScaler().fit_transform(statsrel[features])
         X.shape

Out[77]: (773, 29)

In [78]: y= statsrel['ERA_next_year'].values
         len(y)

Out[78]: 773
```

Then we create a train-test split so that we can accuartely test our regressor later.

```
In [95]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33)
```

Now we can use a linear regressor on the data to try and make predictions

```
In [96]: regressor = LinearRegression()
         regressor.fit(X_train, y_train)

Out[96]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=1, normalize=False)

In [97]: y_prediction = regressor.predict(X_test)
```

We calculate the root mean squared error (RMSE) of our predictions:

```
In [98]: RMSE = np.sqrt(mean_squared_error(y_true = y_test, y_pred = y_prediction))
         RMSE

Out[98]: 0.74148067269441
```

Our linear regressor on average is .74 runs off in its prediction. For comparison, lets look at our errors if we use the pitchers ERA from the previous year as our predicition

```
In [88]: RMSEERA = np.sqrt(mean_squared_error(y_true = statsrel['ERA_next_year'], y_pred = stats
         RMSEERA

Out[88]: 0.8972064594000275
```

Our regressor fares slightly better, how about if we use SIERA and xFIP?

```
In [92]: RMSESIERA = np.sqrt(mean_squared_error(y_true = statsrel['ERA_next_year'], y_pred = sta
         RMSESIERA

Out[92]: 0.753193803454823

In [93]: RMSExFIP = np.sqrt(mean_squared_error(y_true = statsrel['ERA_next_year'], y_pred = stat
         RMSExFIP

Out[93]: 0.7460069642340079
```

Our regressor fares about the same as xFIP and SIERA, which is no surprise given they both largely use stats from our list of features as inputs.

In conclusion, it seems that in terms of raw statistics a pitchers strikeout rates are the most important stats to use in order to predict his ERA, and the ERA estimators xFIP and SIERA both do a good job in taking this data into account.